

Seminar

2025.08.22

Nanxing Hu



DeepEyes: Incentivizing “Thinking with Images” via Reinforcement Learning

**Ziwei Zheng^{1,2*}, Michael Yang^{1*}, Jack Hong^{1*}, Chenxiao Zhao^{1*†},
Guohai Xu^{1‡}, Le Yang^{2‡}, Chao Shen², XingYu¹**

¹Xiaohongshu Inc., ² Xi'an Jiaotong University

* Equal contribution, Random order † Main Code Contributor ‡ Corresponding Author

Project Homepage

{chenxiao2, xuguohai}@xiaohongshu.com, yangle15@xjtu.edu.cn,
ziwei.zheng@stu.xjtu.edu.cn, {yangminghao199, jaaackhong}@gmail.com

Setting Clarification

Setting: (high resolution) image understanding by VLM

Example:



Q: Is the clock to the left of the laptop?

A: ???????

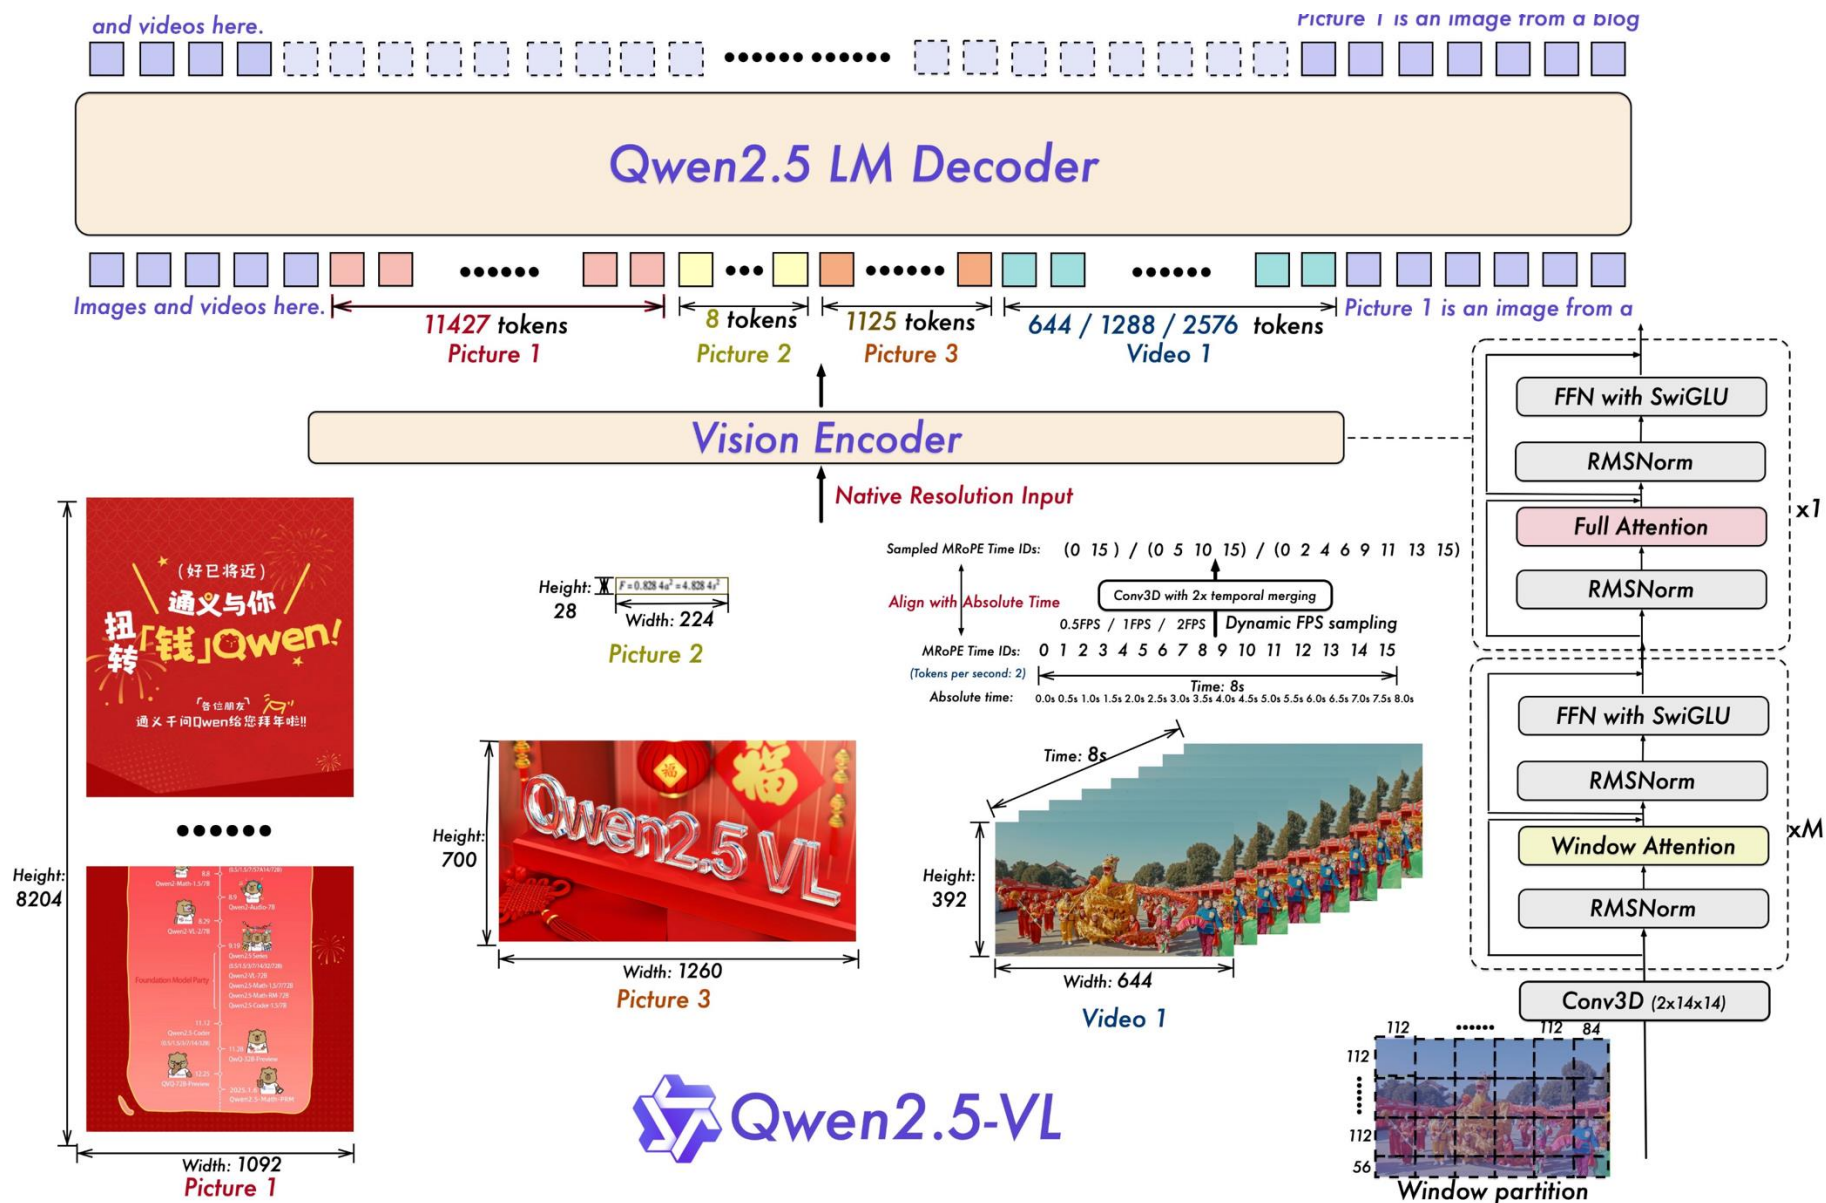


A: Yes !

If we can make VLM recall local visual information in the reasoning process, the reasoning ability will be improved.

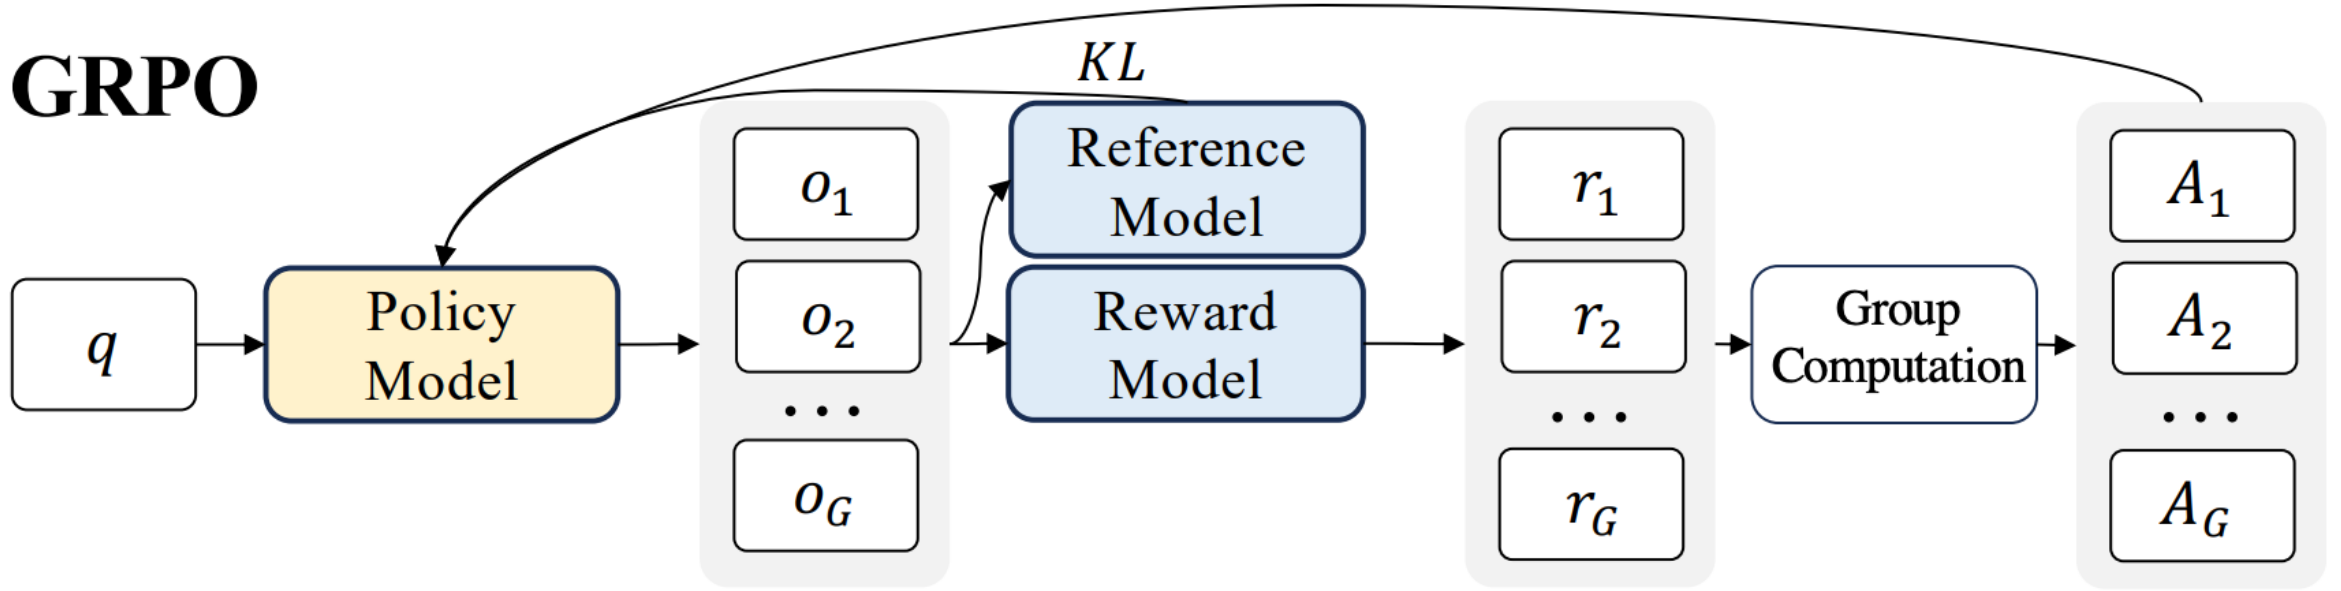
Approach: Reinforcement Learning

Preliminary: VLM



Preliminary: GRPO

GRPO

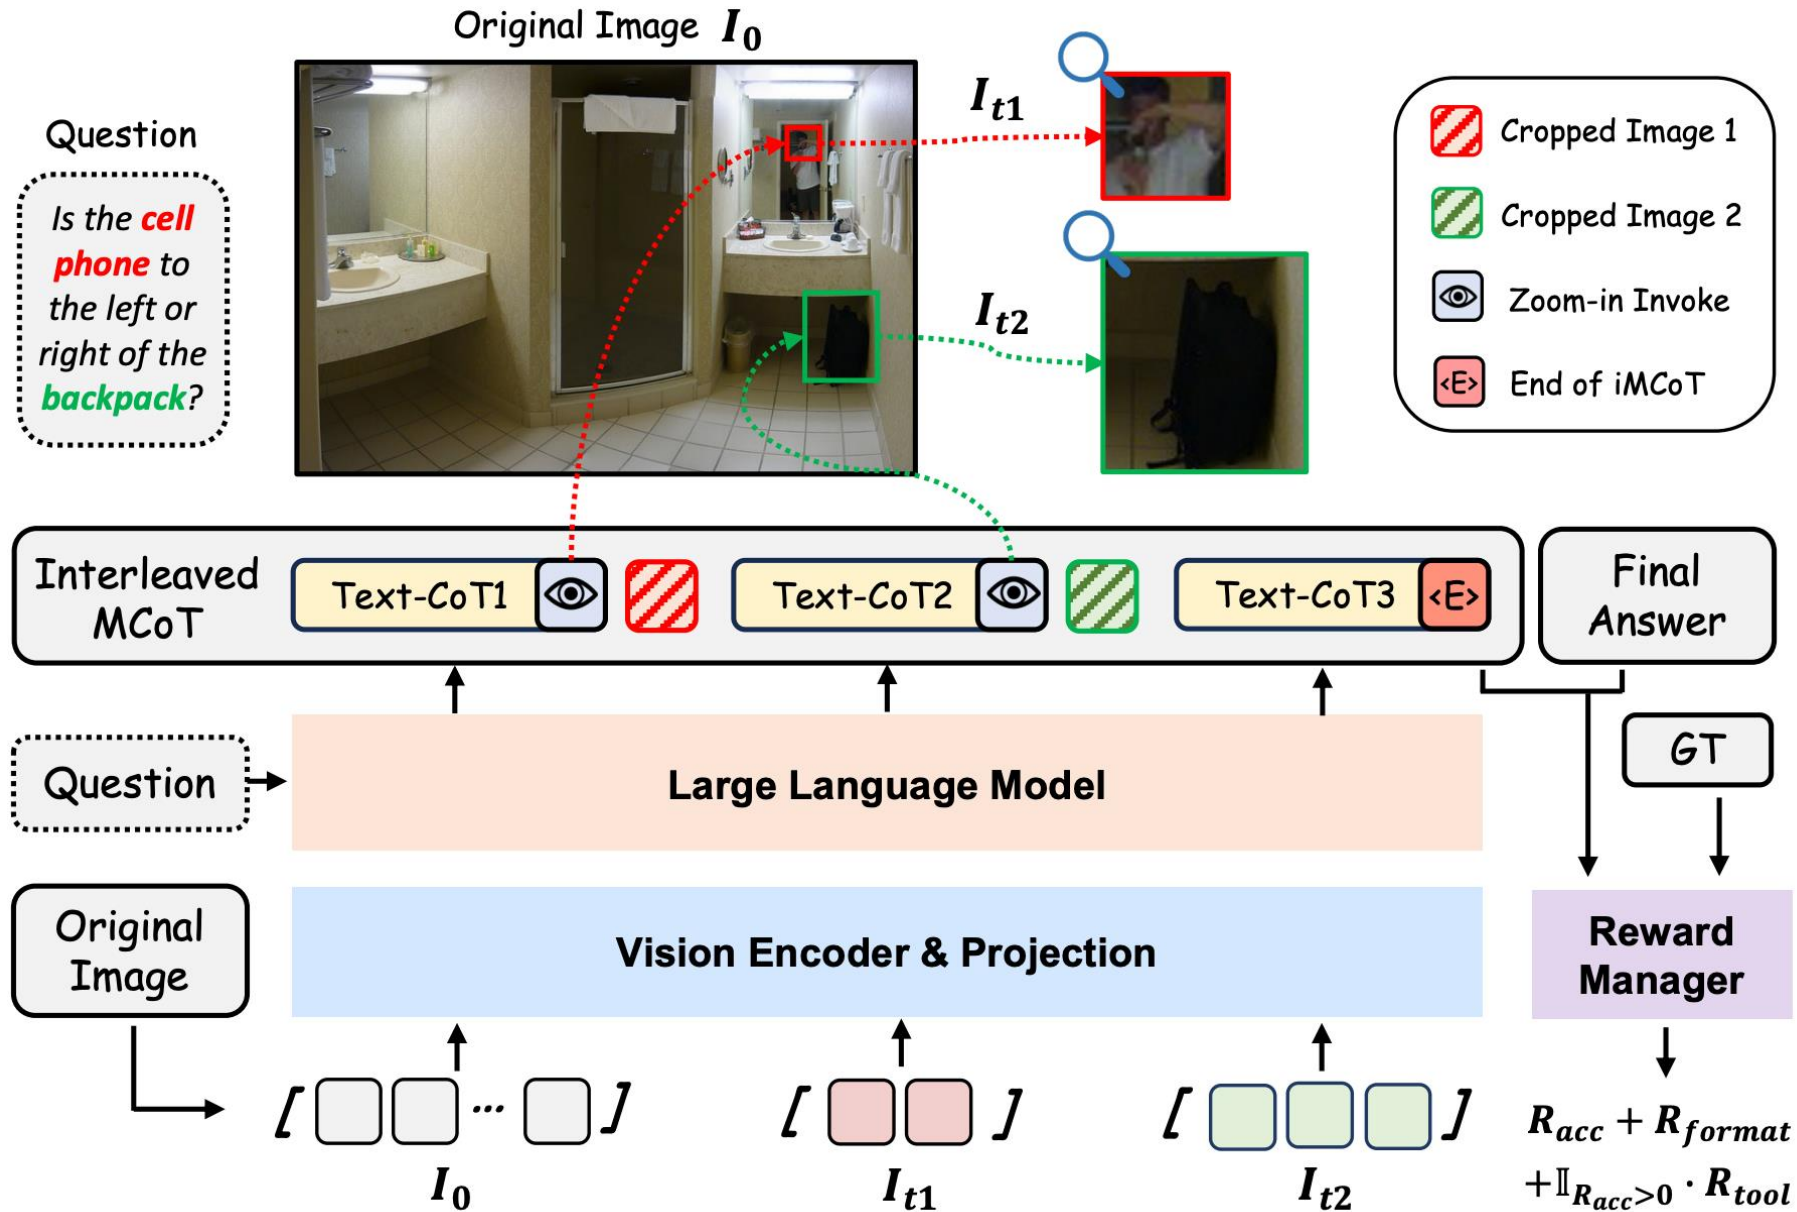


$$\hat{A}_{i,t} = \tilde{r}_i = \frac{r_i - \text{mean}(\mathbf{r})}{\text{std}(\mathbf{r})},$$

$$\mathcal{J}_{GRPO}(\theta) = \mathbb{E}[q \sim P(Q), \{o_i\}_{i=1}^G \sim \pi_{\theta_{old}}(O|q)]$$

$$\frac{1}{G} \sum_{i=1}^G \frac{1}{|o_i|} \sum_{t=1}^{|o_i|} \left\{ \min \left[\frac{\pi_{\theta}(o_{i,t}|q, o_{i,<t})}{\pi_{\theta_{old}}(o_{i,t}|q, o_{i,<t})} \hat{A}_{i,t}, \text{clip} \left(\frac{\pi_{\theta}(o_{i,t}|q, o_{i,<t})}{\pi_{\theta_{old}}(o_{i,t}|q, o_{i,<t})}, 1 - \epsilon, 1 + \epsilon \right) \hat{A}_{i,t} \right] - \beta \mathbb{D}_{KL} [\pi_{\theta} || \pi_{ref}] \right\}, \quad (3)$$


Pipeline




Data Construction

Principle:

- Diverse Tasks and Image Distribution
- Reasoning Ability Enhancement
- Tool Effectiveness



Data Collection:



fine-grained data, chart data, and reason data.

Data Selection:

Managing Difficulties: generate 8 responses per question excluded as they are either too easy or too hard.

+

Facilitating Tool Integration: select instances achieves correct results when utilizing ground-truth crop regions.

Reward design

$$R(\tau) = R_{\text{acc}}(\tau) + R_{\text{format}}(\tau) + \mathbb{I}_{R_{\text{acc}}(\tau) > 0} \cdot R_{\text{tool}}(\tau),$$

$R_{\text{acc}}(\tau)$: accuracy reward assesses whether the final answer is correct.

$R_{\text{format}}(\tau)$: formatting reward penalizes poorly structured outputs.

$R_{\text{tool}}(\tau)$: Tool usage bonus is awarded only when the model produces a correct answer and invokes at least one external perception tool during the trajectory

Experiment results

Model	E2E	Param Size	V* Bench [41]			HR-Bench 4K [59]			HR-Bench 8K [59]		
			Attr	Spatial	Overall	FSP	FCP	Overall	FSP	FCP	Overall
GPT-4o [60]	✓	-	-	-	66.0	70.0	48.0	59.0	62.0	49.0	55.5
o3 [8]	✓	-	-	-	95.7	-	-	-	-	-	-
SEAL [41]	✗	7B	74.8	76.3	75.4	-	-	-	-	-	-
DyFo [44]	✗	7B	80.0	82.9	81.2	-	-	-	-	-	-
ZoomEye [61]	✗	7B	93.9	85.5	90.6	84.3	55.0	69.6	88.5	50.0	69.3
LLaVA-OneVision [62]	✓	7B	75.7	75.0	75.4	72.0	54.0	63.0	67.3	52.3	59.8
Qwen2.5-VL* [58]	✓	7B	73.9	67.1	71.2	85.2	52.2	68.8	78.8	51.8	65.3
Qwen2.5-VL* [58]	✓	32B	87.8	88.1	87.9	89.8	58.0	73.9	84.5	56.3	70.4
DeepEyes	✓	7B	91.3	88.2	90.1	91.3	59.0	75.1	86.8	58.5	72.6
Δ (vs Qwen2.5-VL 7B)	-	-	+17.4	+21.1	+18.9	+6.1	+6.8	+6.3	+10.0	+6.8	+7.3

High-Resolution Benchmarks

Grounding and Hallucination Benchmarks

Model	Param Size	refCOCO	refCOCO+	refCOCOG	ReasonSeg	POPE			
						Adversarial	Popular	Random	Overall
LLaVA-OneVision [62]	7B	-	-	-	-	-	-	-	88.4
Qwen2.5-VL [58]	7B	90.0	84.2	87.2	-	-	-	-	-
Qwen2.5-VL* [58]	7B	89.1	82.6	86.1	68.3	85.9	86.5	87.2	85.9
DeepEyes	7B	89.8	83.6	86.7	68.6	84.0	87.5	91.8	87.7
Δ (vs Qwen2.5-VL 7B)	-	+0.7	+1.0	+0.6	+0.3	-1.9	+1.0	+4.6	+1.8

Model	Param Size	Math Vista [64]	Math Verse [65]	Math Vision [66]	We Math [67]	Dyna Math [68]	Logic Vista [69]
LLaVA-OneVision [62]	7B	58.6 [†]	19.3 [†]	18.3 [†]	20.9 [†]	-	33.3 [†]
Qwen2.5-VL [58]	7B	68.2	49.2	25.1	35.2 [†]	-	44.1 [†]
Qwen2.5-VL* [58]	7B	68.3	45.6	25.6	34.6	53.3	45.9
DeepEyes	7B	70.1	47.3	26.6	38.9	55.0	47.7
Δ (vs Qwen2.5-VL 7B)	-	+1.9	+1.7	+1.0	+4.3	+1.7	+1.8

Multimodal Reasoning Benchmarks

Ablations

Table 4: **Ablation Study on iMCoT.** We compare the results of RL training using text-only CoT and iMCoT on the same datasets.

Model	Attr	V* Bench		HR-Bench 4K			HR-Bench 8K		
		Spatial	Overall	FSP	FCP	Overall	FSP	FCP	Overall
Qwen2.5-VL [58]	73.9	67.1	71.2	85.2	52.2	68.8	78.8	51.8	65.3
RL w. Text-only CoT	90.4	85.5	88.5	92.3	58.5	75.4	69.3	52.3	60.8
DeepEyes	91.3	88.2	90.1	91.3	59.0	75.1	86.8	58.5	72.6

Multi-modal CoT make a difference

Table 5: **Impact of Training Data.** Fine represents the fine-grained data. HR denotes HR-Bench. Row #0 is the origin score of Qwen2.5-VL 7B.

#	Fine	Reason	Chart	High-Resolution			Basic VL Capability		Reasoning	
				V* Bench	HR-4K	HR-8K	ReasonSeg	POPE	MathVista	MathVerse
0				71.2	68.8	65.3	68.3	85.9	68.2	45.6
1	✓			86.9	68.9	67.3	69.0	86.6	67.0	42.9
2	✓			91.6	74.1	71.0	69.1	88.1	64.7	41.3
3	✓	✓		91.6	73.8	70.5	68.6	88.8	67.7	43.8
4	✓		✓	90.1	74.6	74.6	68.5	87.9	64.6	38.1
5	✓	✓	✓	90.1	75.1	72.6	68.6	87.7	70.1	47.3

#1 denote training with unfiltered data

- data selection is necessary
- reasoning data is necessary for maintain reasoning ability
- Chart data can benefit the Math problem

Reinforcement Learning Tuning for VideoLLMs: Reward Design and Data Efficiency

Hongyu Li^{1*}, Songhao Han^{1*}, Yue Liao^{2*†}, Junfeng Luo³, Jialin Gao³,
Shuicheng Yan², Si Liu^{1‡}

¹ BUAA ² NUS ³ Meituan

Setting Clarification

Setting: Reinforcement learning for video-specific reasoning capabilities of MLLMs



Discrete Reward in VideoQA

What are these people chasing in these scene transitions?

- ❌ (A) The man inside the car
- ✅ (B) A drone in the sky
- ❌ (C) A woman on the road
- ❌ (D) A tree in the grass

<think> The subjects travel from a paved highway ... As they move into open grassland, a drone appears overhead. At a lakeside, ... continuous human-computer interaction for aerial surveillance. *</think>*

Video QA

Continuous Reward in Temporal Grounding

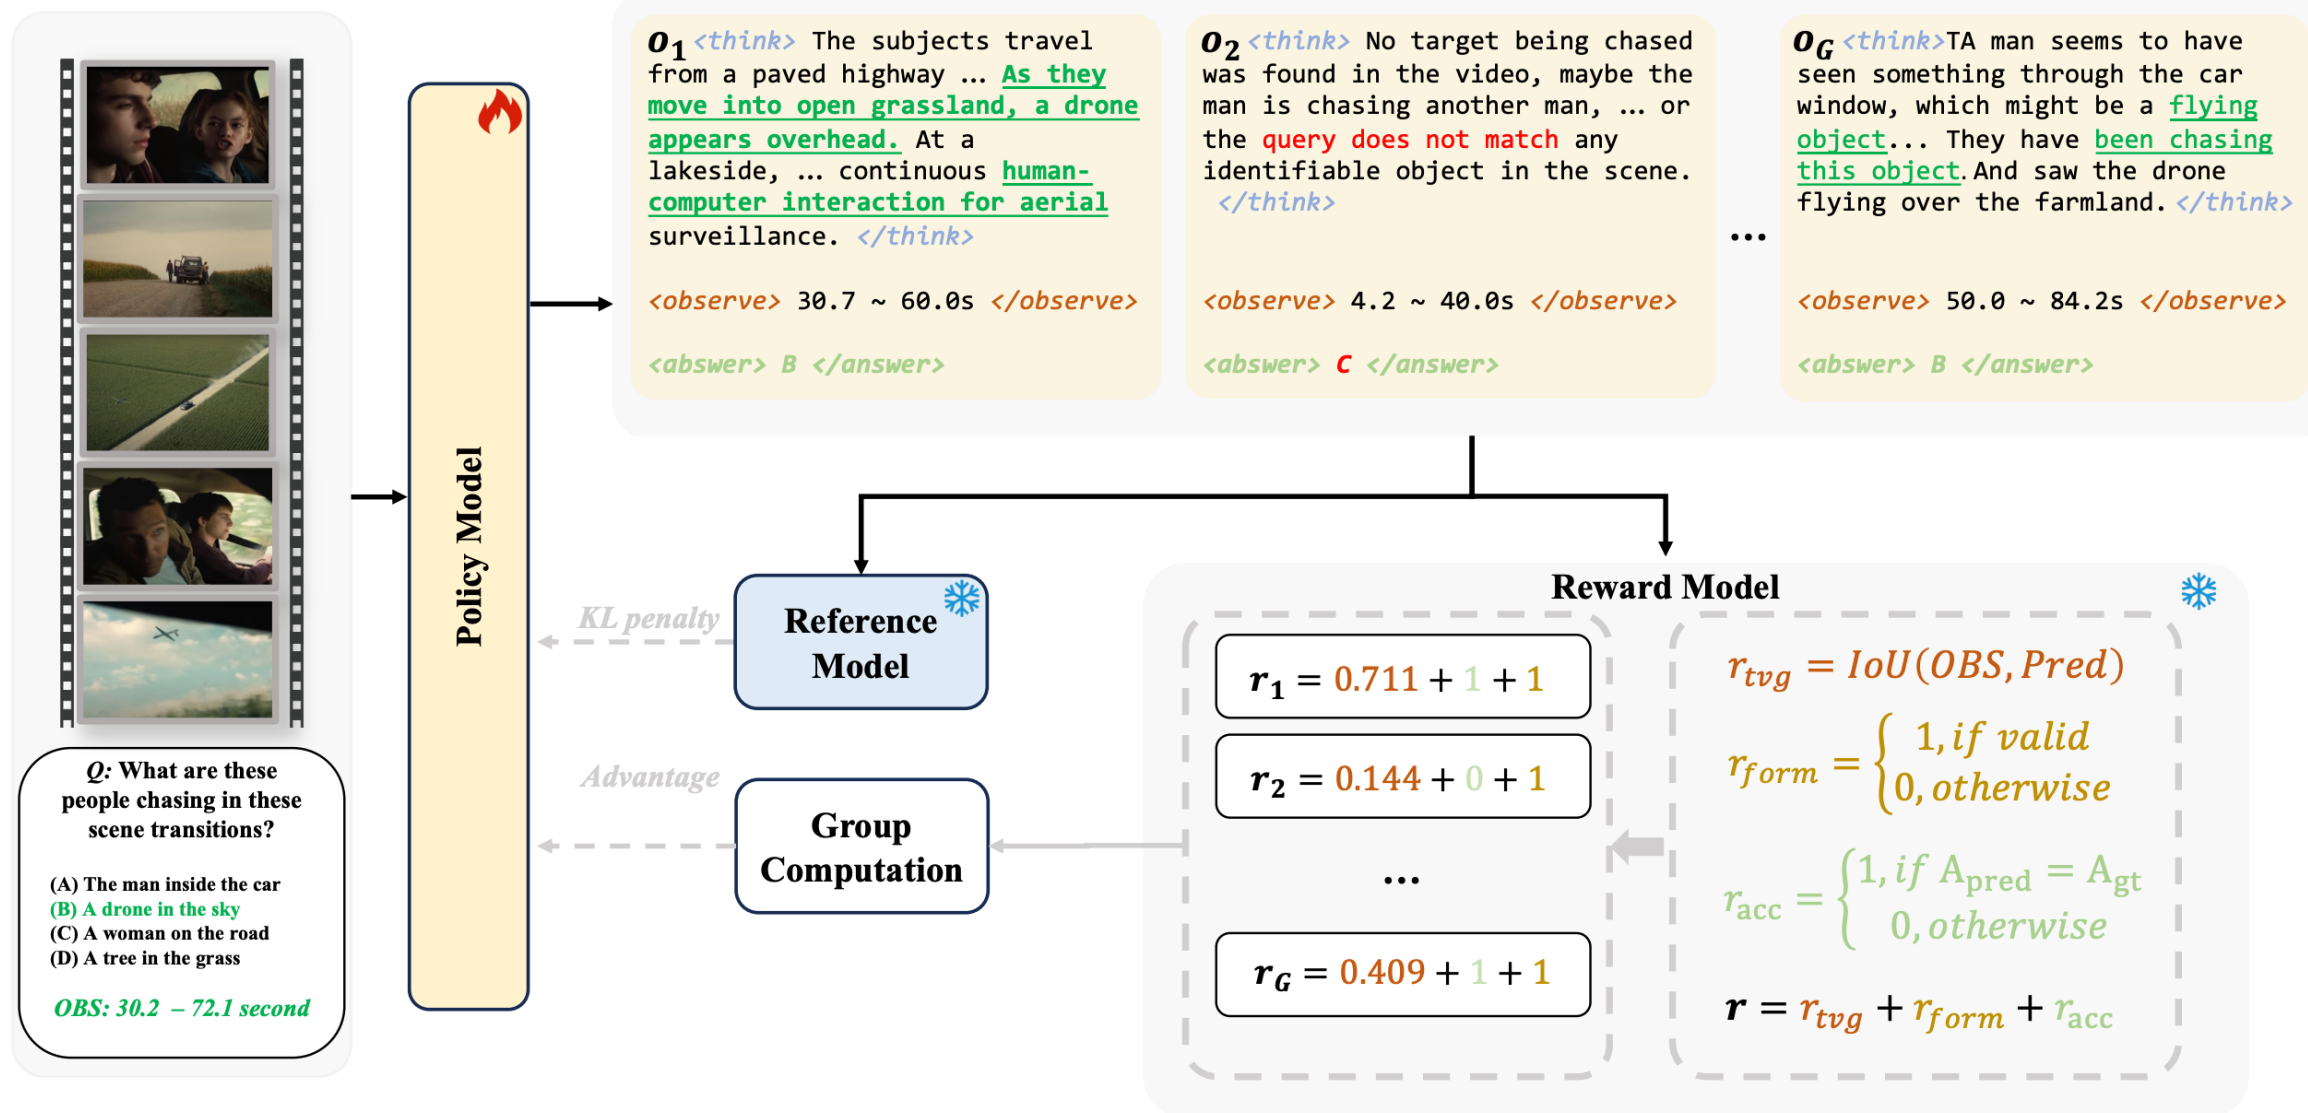
When does the vehicle drive in the cornfield?



<think> The subjects depart from a rural road and drive into expansive cornfields, ... As the vehicle speeds through the crops, an aircraft ... The vehicle eventually emerges into open terrain ... *</think>*


Temporal Grounding

Pipeline

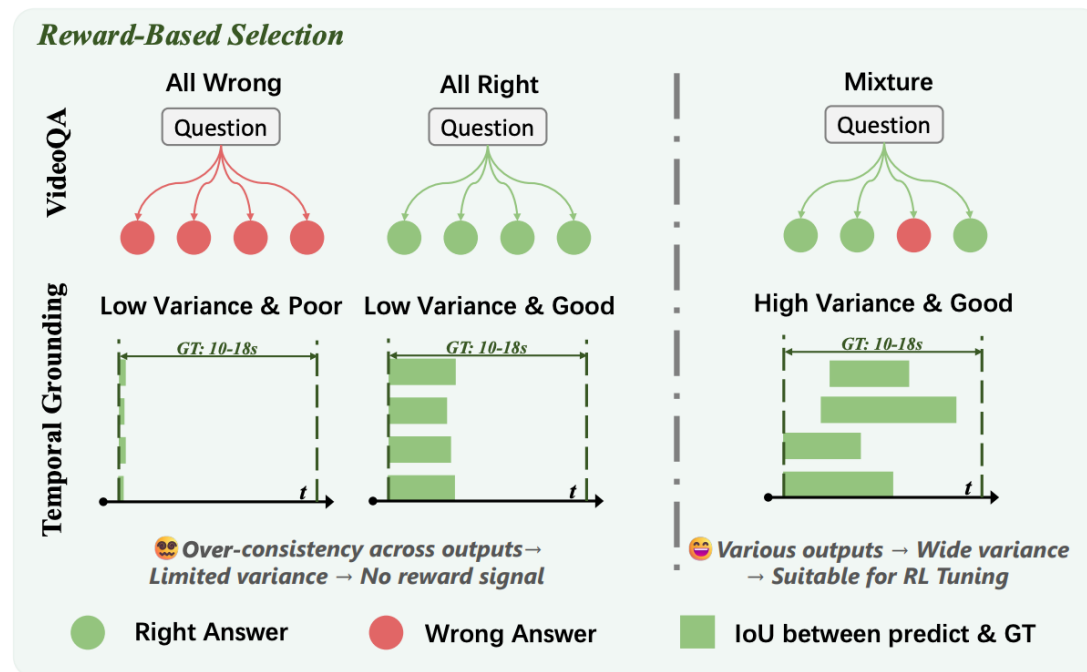


Data Construction

Data construction:



Data Selection :



VQA:

Easy if $c \geq \tau_{\text{easy}}$, Hard if $c \leq \tau_{\text{hard}}$, otherwise Medium,

$$\tau_{\text{easy}} = 1; \tau_{\text{hard}} = 7$$

VTG:

$$\Delta_{\text{IoU}} = \max_i \text{IoU}_i - \text{mean}_i(\text{IoU}_i), \quad \Delta_{\text{IoU}} = 0.3$$

Reward design

Multi-Choice VideoQA (MC-QA) :

$$R_{\text{mc}} = R_{\text{format}} + R_{\text{acc}},$$

Temporal Video Grounding (TVG) :

$$R_{\text{tvg}} = R_{\text{format}} + R_{\text{IoU}},$$

Grounded VideoQA (GQA) :

$$R_{\text{gqa}} = R_{\text{format}} + \frac{1}{2}(R_{\text{acc}} + R_{\text{IoU}}),$$

Experiment results

Method	Temporal Video Grounding			General VideoQA			Reasoning QA	Grounded QA	
	Charades	ANet	ANet-RTL	MVBench	TempCompass	VideoMME	MMVU	NextGQA	
	mIoU	mIoU	mIoU	Avg	Avg	Avg (wo sub)	Avg	mIoU	acc
General VideoLLM									
LLaMA-VID[18]	-	-	-	41.9	45.6	-	-	-	-
VideoLLaMA2[3]	-	-	-	54.6	-	47.9	44.8	-	-
LongVA-7B[39]	-	-	-	-	56.9	52.6	-	-	-
Video-UTR-7B[35]	-	-	-	58.8	59.7	52.6	-	-	-
LLaVA-OV-7B[14]	-	-	-	56.7	-	58.2	49.2	-	-
Kangeroo-7B[19]	-	-	-	61.1	62.5	56.0	-	-	-
GRPO-based Method and Baseline									
Qwen-VL-2.5[2]	28.0	24.0	6.0	65.3	70.9	56.1	61.3	20.2	77.2
Qwen-VL-2.5-SFT	43.0	24.3	18.1	62.0	68.7	49.6	52.5	28.3	70.6
Video-R1[4]	-	-	-	62.7	72.6	57.4	64.2	-	-
Temporal-RLT (ours)	57.0	39.0	27.6	68.1	73.3	57.6	65.0	37.3	78.7

Ablations

Table 4: Ablation Studies: Video QA and TVG Data Selection.

Easy: Middle: Hard	General VideoQA			Reasoning QA	Δ_{IoU}	Charades-STA			
	MVBench	TempCompass	VideoMME	MMVU		Recall@0.3	Recall@0.5	Recall@0.7	mIoU
4 : 4 : 2	64.3	70.0	52.8	59.5	0	78.2	63.9	37.4	54.7
2 : 4 : 4	65.9	70.3	55.9	63.0	0.1	78.0	64.8	38.9	54.9
2 : 6 : 2	67.2	71.3	56.8	62.1	0.2	78.8	63.7	38.5	55.0
1 : 8 : 1	68.1	73.4	57.1	63.4	0.3	78.6	64.5	39.9	55.5
0 : 10 : 0	68.1	72.5	58.6	63.1					

(a) Ablation for Video QA Data Selection

(b) Ablation for TVG Data Selection

Diversity and training efficacy data makes a difference.

Table 5: Temporal Video Grounding OOD Evaluation.

Tuning Type	Charades-STA				ActivityNet				ActivityNet-RTL			
	R@0.3	R@0.5	R@0.7	mIoU	R@0.3	R@0.5	R@0.7	mIoU	R@0.3	R@0.5	R@0.7	mIoU
X	42.4	29.8	14.0	28.0	34.4	22.5	11.6	24.0	7.9	2.6	2.9	6.0
SFT	73.9	61.6	38.5	52.8	33.4	18.9	9.0	23.1	24.0	14.8	7.4	17.8
RLT	80.2	68.3	44.5	57.9	56.9	38.4	20.2	39.1	40.2	22.7	10.9	26.3

only trained on Charades-STA dataset

RLT performs significantly better than SFT on OOD task.